

OPERATING SYSTEMS CSCI-SHU215

Lecture 04 - Processes

Notion of Process

2

Definition (C. Girault)

A process is the sequential execution of the instructions of a program in a **Namespace**

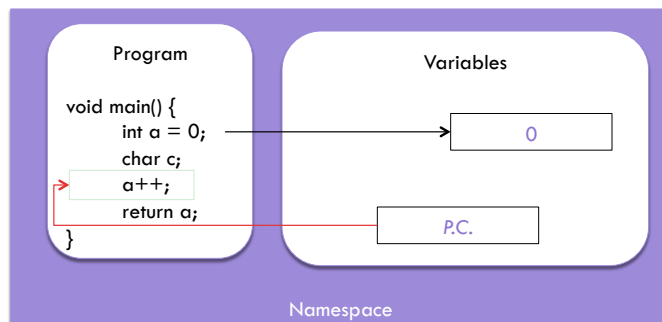
Namespace

- ▣ code
- ▣ stack
- ▣ shared variables and constants
- ▣ files

Multiple runs of the same program in different namespaces produce different processes

Notion of Process

3



One register, the **program counter**, stores the current running instruction

Notion of Process


4

A process is an active entity of the system

- ▣ Corresponds to the execution of a binary program
- ▣ Identified in a unique way by its **pid** number
- ▣ Has 3 segments: code, data, and stack
- ▣ Runs under the identity of a user
- ▣ Has a current directory

Notion of Process

5



Each process is independent

Two processes can be associated with the same program (code)


Program Execution

6

```

1:  int cont;
2:  void foo (int max) {
3:      int i;
4:      for (i=0; i++; i<max)
5:          printf ("%d \n", i);
6:  }

7:  int main (int argc, char* argv []) {
8:      int cont=5;
9:      foo(cont) ;
10:     return EXIT_SUCCESS;
11:  }
    
```



Data Segment

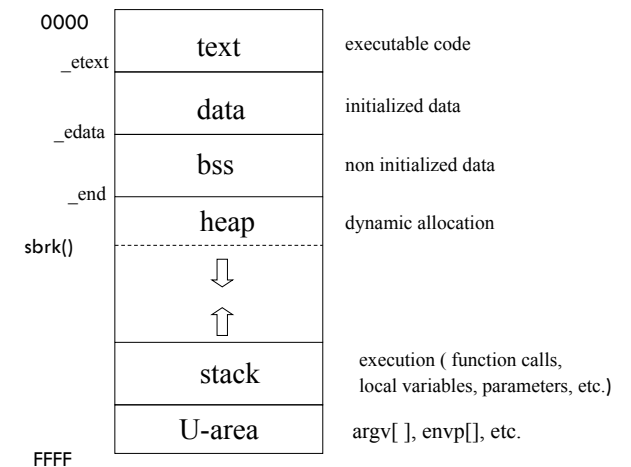
7

- **Data**
Data initialized upon loading the process
- **BSS (Block Started by Symbol)**
Non initialized data
- **Heap**
Dynamically allocated memory areas
eg. malloc(), calloc()

The heap and the stack grow in opposite directions

Namespace of a process

8



Dynamic allocation of memory

9

```
#include <stdlib.h>
```

```
void * malloc(size_t size)  
    allocates a block of size bytes
```

```
void * calloc(size_t nb, size_t size)  
    allocates a block of nb * size bytes (initially set to 0)
```

```
void * realloc(void* ptr, size_t size)  
    resizes a block of memory (preserves content)
```

```
void free (void * ptr)  
    frees the allocated memory
```

States of a process

10

At runtime, a process switches states

Running

process instructions are being executed

Suspended

process is waiting for a resource to become available

Ready

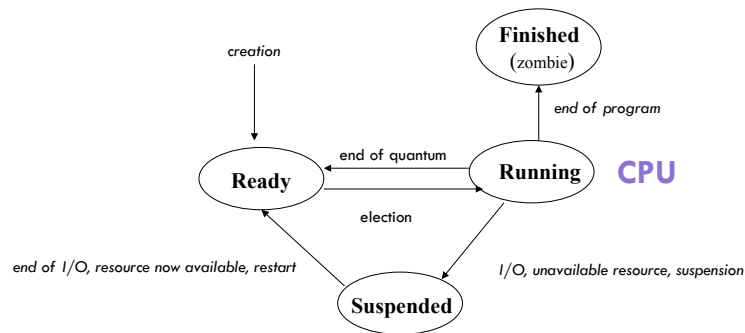
process waits to be assigned to a processor

Zombie

process has finished running, but its parent has yet to acknowledge its termination

States of a process

11



Quantum: unit of duration (e.g. 100 ms)

Attributes of a process

12

Identity of a process

pid: positive integer (type POSIX `pid_t`)

```
pid_t getpid (void)  
    returns the pid of a process
```

Exemple:

```
#define _XOPEN_SOURCE 700  
#include <sys/types.h>  
#include <unistd.h>  
#include <stdio.h>  
  
int main(int argc, char **argv) {  
    printf (" process pid is: %d \n", getpid()) ;  
    return EXIT_SUCCESS;  
}
```

Attributes of a process

13

A process is tied to a user and its group

UID (User identifier) & GID (group identifier)

Identities associated with access permissions by the kernel

Permissions

Rights granted to the user (group) that starts the program

Effective rights

Rights granted to the program itself

- identity that the kernel takes into account to check the access permissions for operations requiring an identification
- eg. opening a file, making a system call

Process Control Block (PCB)

14

Information associated with each process

- State of the process
- ID of the process
- Program Counter (PC)
- CPU registers
- CPU scheduling info
- Memory management info
- Accounting info
- I/O info

15

Process creation with fork

Fork - creation of a process

16

Primitive *pid_t* *fork* (*void*)

Dynamic creation of a new process (*child*)

The child is executed concurrently with the process that created it (*parent*)


```
#include <sys/types.h>
#include <unistd.h>
pid_t fork (void)
```

Created child process is a copy of the parent process

Fork - creation of a process

17

- Both processes share the same physical code
- Duplication of the stack and data segment
 - variables of the child have the same values as those of the parent upon *fork* call;
 - any value modification of a variable by one of the processes is not visible to the other



Fork - creation of a process

18

One single call, but two return values

Each process resumes its execution

Return values differ

0	returned to the child
<i>child process pid</i>	returned to the parent
-1	system call failed
errno <errno.h>:	
ENOMEM	not enough memory available
EAGAIN	too many processes created

`pid_t getppid (void)`
returns the pid of the parent

Fork - creation of a process

19

Example 1

```
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

int main (int argc, char* argv []) {
    pid_t pid_child;
    switch (pid_child = fork ( ) ) {
        case (pid_t) -1:
            perror ("fork"); exit (1);
        case (pid_t) 0:
            printf ("CHILD> pid %d, parent pid %d \n", getpid (), getppid ( ) );
            return EXIT_SUCCESS;
        default:
            printf ("PARENT> pid %d, child pid %d \n", getpid(), pid_child);
            return EXIT_SUCCESS;
    }
}
```

test-fork1.c

Fork - creation of a process

20

Concurrency

- Child gets to run before its parent

```
$test-fork1
CHILD> pid 1528, parent pid 1476
PARENT> pid 1476, child pid 1528
```
- Parent gets to run before its child

```
$test-fork1
PARENT> pid 1476, child pid 1528
CHILD> pid 1528, parent pid 1
    Child becomes an orphan, gets adopted by init process (pid = 1)
```

Fork - data duplication

21

Example 2:

```
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

int main (int argc, char* argv []) {
    int a= 3; pid_t pid_child;
    a *=2;
    if ( (pid_child = fork ( ) ) == -1 ) {
        perror ("fork"); exit (1);
    } else
        if (pid_child == 0) {
            a=a+3;
            printf ("child a=%d \n", a);
        } else
            printf ("parent a=%d \n", a);
    return EXIT_SUCCESS;
}
```

test-fork2.c

```
$test-fork2
child a=9
parent a=6

$test-fork2
parent a=6
child a=9
```

Fork - looped calls


22

Example 3: a process creates N childs

test-fork3.c

```
#include <sys/types.h>
#include <unistd.h>
#define N 3

int main (int argc, char* argv []) {
    int i=0; pid_t pid;
    while (i <N) {
        if (((pid=fork ( ) )== -1) || (pid==0))
            exit (1);
        i++;
    }
    return EXIT_SUCCESS;
}
```




Fork - looped calls

23

Example 4: How many processes at the end?

test-fork4.c



```
#define _XOPEN_SOURCE 700
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

int main (int argc, char* argv []) {
    int i =0 ;
    while (i <3) {
        printf ( "%d ", i);
        i ++;
        if (fork ( ) == -1)
            exit (1);
    }
    printf( "\n ");
    return EXIT_SUCCESS;
}
```

Fork - Inheritance

24

A child process inherits

- ▣ User ID and group ID (real and effective)
- ▣ Session ID
- ▣ Current working directory
- ▣ Current umask bits, signal mask, signal routines
- ▣ Attached shared memory
- ▣ Environment variables
- ▣ Open file descriptors
- ▣ **nice** value
- ▣ ...

Fork - Inheritance

25

A child process does not inherit

- ❑ Identity (pid) of its parent process
- ❑ Execution time
- ❑ Pending signals
- ❑ File locks held by the parent
- ❑ Alarms and timers

functions alarm, setitimer,...

Fork - Inheritance vs dynamic memory

26

test fork5.c


```
#DEFINE SIZE 100
char* ptr1=NULL;
char* ptr2 =NULL;
int main (int argc, char* argv []) {
    pid_t pid;
    ptr1= (char*)malloc(SIZE);
    memcpy(ptr1, (char*)"toto", strlen("toto"));
    if ((pid=fork ( )) == -1)
        exit (1);
    else
        if (pid != 0){
            /*parent */
            ptr2= (char*) malloc(SIZE);
            memcpy(ptr2, (char*) "titi", strlen("titi"));
            printf ("PARENT - ptr1:%s ; ptr2:%s \n", ptr1, ptr2);
        }
        else
            printf ("CHILD - ptr1:%s ; ptr2:%s \n", ptr1, ptr2);

    return 0;
}
```

```
$test_fork5
PARENT - ptr1: toto; ptr2:titi
CHILD - ptr1: toto; ptr2:
```

Fork - Inheritance vs file descriptors

27



28

End of a process

Process termination

29

`exit(int val)` or `return val`

`val` is a value that the parent can acquire

Pre-defined constants

`EXIT_SUCCESS`

`EXIT_FAILURE`

If a process started by the shell terminates with an error

error code available in `variable $?`

`echo $?`

Process termination

30

When a process terminates correctly

All I/O streams get closed

Buffers are emptied

Call to `_exit()` (system call)

closes open file descriptors

parent process received a signal `SIGCHLD`

The terminated process becomes a *zombie*

Process termination

31

Zombie process

Temporary state of a process until its parent acknowledges its termination

Parent/Child Synchronization

Upon termination, either with an `exit` call or with a `return` instruction in the main function, a process assigns a value to its `return code`

The parent process can access this value by calling function `wait/waitpid`

32

Basic parent/child synchronization

Basic parent/child synchronization

33

Primitive `pid_t` `wait` (`int*` `status`)

```
#include <sys/types.h>
#include <sys/wait.h>
pid_t wait (int* status)
```

If the calling process

- has at least one zombie child
 - call returns the identity of one of the zombies
 - recovery of the return code value in variable `status`
- has a child, but none is in a zombie state
 - process is blocked until one of its children becomes zombie
 - notification by a signal `SIGCHLD`
- does not have any children
 - the call returns -1 and the value of `errno` becomes `ECHILD`

Basic parent/child synchronization

34

Interpretation of the return value `int *status`

Pre-defined **macros** to solve portability issues

Type of termination

`WIFEXITED`: not NULL if the child process terminated normally
`WIFSIGNALED`: not NULL if the child process terminated because of a signal
`WIFSTOPPED`: not NULL if the child process is stopped (option `WUNTRACED` `waitpid`)

Information on the return value or the signal

`WEXITSTATUS`: return code if the process ended normally
`WTERMSIG`: value of the signal that caused the process termination
`WSTOPSIG`: value of signal that stopped the process

Basic parent/child synchronization

35

include files: `stdio.h`, `sys/types.h`, `sys/wait.h`, `unistd.h`, `stdlib.h`

test-wait.c

```
int main(int argc, char **argv) {
    pid_t pid_child; int status;
    if (fork () == 0) {
        printf ("CHILD:  pid = %d \n", getpid());
        exit (2);
    } else {
        pid_child = wait(&status);
        if (WIFEXITED (status) ) {
            printf("PARENT: child %d ended with status %d \n",
                pid_child, WEXITSTATUS (status));
            return EXIT_SUCCESS;
        } else
            return EXIT_FAILURE;
    }
}
```

```
$test-wait
CHILD: pid = 3254
PARENT: child 3254 ended with status 2
```

Basic parent/child synchronization

36

include files: `stdio.h`, `sys/types.h`, `sys/wait.h`, `unistd.h`, `stdlib.h`

test-wait2.c

```
int main(int argc, char **argv) {
    pid_t pid_child; int status;
    if (fork () == 0) {
        printf ("CHILD:  pid = %d \n", getpid());
        pause ();
        exit (2);
    } else {
        pid_child = wait(&status);
        if (WIFEXITED (status) ) {
            printf ( "PARENT: child %d ends with status %d \n",
                pid_child, WEXITSTATUS (status));
            return EXIT_SUCCESS;
        } else
            if (WIFSIGNALED (status) ) {
                printf ( "PARENT: child %d ended by signal %d \n",
                    pid_child, WTERMSIG (status));
                return EXIT_SUCCESS;
            }
        return EXIT_FAILURE;
    }
}
```

```
$test-wait2 &
CHILD: pid= 4897
$kill -KILL 4987
PARENT: child 4987 ended by signal 9
```

Basic parent/child synchronization

37

```
#define N 3
int cont = 0;
int main (int argc, char* argv []) {
    int i=0; pid_t pid;
    while (i < N) {
        if ((pid=fork ( ) )== 0) {
            cont++;
            break;
        }
        i++;
    }

    if (pid != 0) {
        /* parent */
        for (i=0; i<N; i++)
            wait (NULL);
        printf ("cont:%d \n", cont);
    }
    return EXIT_SUCCESS;
}
```

test-wait3.c

What is the value
displayed for *cont* ?

Basic parent/child synchronization

38

- Primitive *pid_t* *waitpid* (*pid_t* *pid*, *int** *status*, *int* *opt*)

```
#include <sys/types.h>
#include <sys/wait.h>
pid_t waitpid (pid_t pid, int* status, int opt )
```

Option *opt* of function *waitpid* allows a parent process to check the termination of a child

- with a specific pid value
- or that belongs to a group |pid|

Basic parent/child synchronization

39

- Value of parameter *pid*
 - > 0 PID of child process
 - 0 any process that belongs to the same group as the caller
 - 1 any child process
 - < -1 any child process in group |pid|
- Value of parameter *opt*
 - WNOHANG non blocking call
 - WUNTRACED check for stopped processes
- Return code
 - -1 error
 - 0 (non blocking mode) process is still running
 - *pid* of the zombie process

40

Code replacement

Code replacement

41

Primitive exec replaces the code that is executed by a new program

Arguments: name of the new program + parameters

The new program will be executed in the namespace of the calling process

If the call succeeds, the new program neverhands control back to the calling process

Examples of error (*errno*)

EACCES	no permission to access the file
ENOENT	file not found...

Code replacement

42

□ argv in list format

```
int execl (const char *path, const char *arg, ...);
```

□ argv in array format

```
int execv (const char *path, char * const argv[]);
```

Last argument value *must always be* NULL

Code replacement

43

Example - execl

```
#include <sys/types.h>
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char **argv) {
    execl ("/usr/bin/wc", "wc", "-w", "/tmp/fichier1", NULL);
    perror ("execl");
    return EXIT_SUCCESS;
}
```

Code replacement

44

Example - execlp

```
#include <sys/types.h>
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char **argv) {
    execlp ("wc", "wc", "-w", "/tmp/fichier1", NULL);
    perror ("execlp");
    return EXIT_SUCCESS;
}
```

Code replacement

45

Example - execv

```
#include <sys/types.h>
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char **argv) {
    char *arg_vect [4];
    arg_vect[0] = "wc";
    arg_vect[1] = "-w";
    arg_vect[2] = "/tmp/fichier1";
    arg_vect[3] = NULL;

    execv ("/usr/bin/wc",arg_vect);
    perror ("execv");
    return EXIT_SUCCESS;
}
```

Code replacement

46

include files: stdio.h, sys/types.h, stdio.h, unistd.h, stdlib.h

```
sleep_exec.c
int main(int argc, char **argv) {
    char* argv_sleep[] = {"sleep_prog", "2", (char *) NULL };
    printf ("Begin - pid:%d \n", getpid ());
    execv ("./sleep_prog", argv_sleep);
    printf("End - pid:%d \n", getpid ());
    return EXIT_SUCCESS;
}
```

```
sleep_prog.c
int main(int argc, char **argv) {
    printf("sleep_prog> begin - pid:%d\n ", getpid ());
    sleep (atoi(argv[1]));
    printf("sleep_prog> end - pid:%d\n ", getpid ());
    return EXIT_SUCCESS;
}
```

```
$ sleep_exec
Begin - pid: 356
sleep_prog> begin - pid: 356
sleep_prog> end - pid:356
```