

OPERATING SYSTEMS CSCI-SHU215

Lecture 06 – Scheduling

2

Fundamental Concepts

Running a Program

3

1. Program = sequence of instructions
2. Task = [1] + data
3. Process = [2] + execution context (PCB)

Monoprogramming

Process retains the processor until its termination

Sharing the Processor

4

CPU-I/O Cycle

The lifecycle of a process consists of a sequence of cycles
[CPU execution + I/O wait]

Distribution of the CPU cycles

Multiprogramming + DMA => Maximum CPU usage


Multiprogramming

5

Batch processing

The active process hands over control:

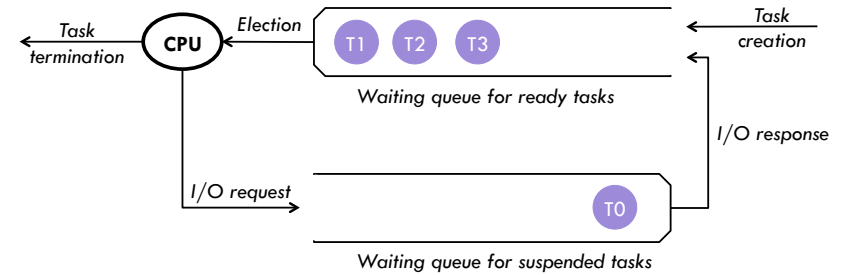
- upon termination
- upon requesting an I/O



Multiprogramming

6

Batch processing scheduler




Multiprogramming

7

Time-sharing

The active process hands over control:

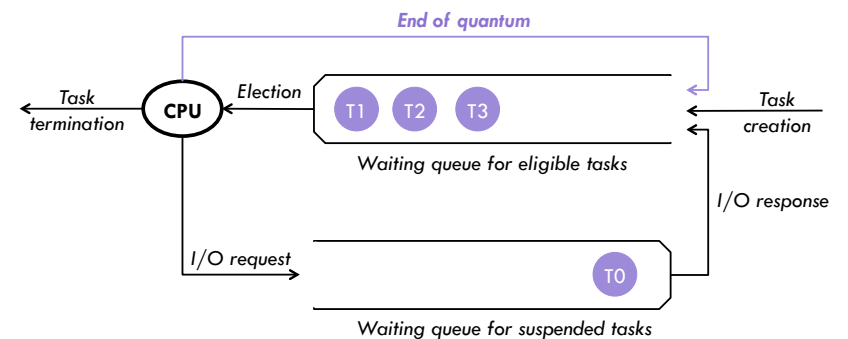
- upon termination
- upon requesting an I/O
- upon spending its quantum



Multiprogramming

8


Time-sharing scheduler



CPU Scheduler

9

- Elects one of the processes that are ready in memory
Allocates the CPU to one of them
- Triggers of an election



Dispatcher

10

Hands over CPU to the elected process

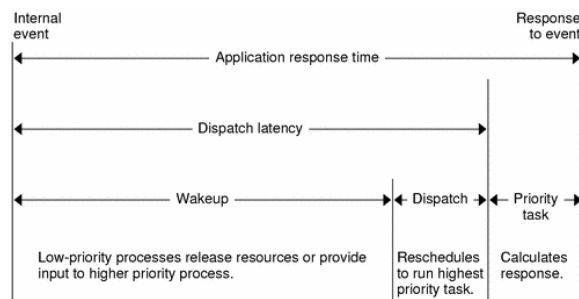
1. Context switch
2. Shift to user mode
3. Restart execution from PC

Dispatch induces overhead

Dispatch Latency

11

Dispatch latency describes the amount of time a system takes to respond to a request for a process to begin operation.



© Oracle

Scheduling Criteria

12

- CPU utilization
- Throughput
of processes completed per unit of time
- Turnaround time
Time required for a particular process to complete, from submission time to completion (wall clock time)
- Waiting time
Time spent by a process in the Ready queue
- Response time
Time taken in an interactive program from the issuance of a command to the commence of a response to that command

Scheduling Policy Objectives

13

Aim towards optimal satisfaction of criteria

In particular:

Guarantee fairness (not equality!)
prevent famines

Many algorithms

eg. FIFO, SJF, RR, EDF, RMS, ...

Choosing the right algorithm depends on system usage
(number of tasks, types, insertion rate, ...)

Optimization Criteria

14

- Maximum CPU utilization
- Maximum throughput
- Minimum turnaround time
- Minimum waiting time
- Minimum response time


15

Common Scheduling Policies

First Come, First Served (FCFS ou FIFO)

16

Process	CPU time	$t_{\text{insertion}}$ (order)	T_{response}	T_{wait}
P_1	24	0(1)	24	0
P_2	3	0(2)	27	24
P_3	3	0(3)	30	27



- Response time (T_{response}) = $t_{\text{end}} - t_{\text{insertion}}$
- Waiting time = $[T_{\text{response}}] - [\text{Computation time}]$

Average waiting time: $(0 + 24 + 27)/3 = 17$ **Convoy Effect!**

FCFS Scheduling

17

Let's change the order of insertions

Process	Comp. Time	$t_{\text{insertion}}$ (order)	T_{response}	T_{wait}
P_1	24	0(3)	30	6
P_2	3	0(1)	3	0
P_3	3	0(2)	6	3



Average waiting time: $(6 + 0 + 3)/3 = 3$

Shortest Job First (SJF)

18

□ Principle

Predict CPU time of each process

Elect the process with the shortest CPU time

□ Two strategies

■ Non preemptive

Once a process obtains the CPU, it will only release it upon termination

■ Preemptive, or *Shortest-Remaining-Time-First (SRTF)*

Upon insertion of a new process, switch to the latter if its CPU time is shorter than that of the running process

□ SJF is optimal


Guarantees that average waiting time is minimal

Example of Non-Preemptive SJF

19

Process	Insertion Time	CPU Time
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

□ SJF (non preemptive)




□ Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$

Example of Preemptive SJF

20

Process	Insertion Time	CPU Time
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

□ SRTF (preemptive)



□ Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$

Predicting the Next Computation Time

21


One simple, fast, and relatively accurate method is the exponential average

t_n duration of the n^{th} CPU burst
 τ_{n+1} prediction for the duration of the next burst
 α weighting factor ($0 \leq \alpha \leq 1$)

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$

Predicting the Next Computation Time

22



CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...


Priority Scheduling

23

- A **priority** (integer > 0) gets assigned to each process
- Scheduler allocates the CPU to the process with the highest priority
 - ▣ Preemptive
 - ▣ Non preemptive
- SJF is a priority-oriented algorithm
Priority value is the expected CPU time
- Problem: Famine
Low priority processes may never run
- Solution: Graceful Ageing
Priority of processes in ready queue increases periodically

Priority Scheduling

24



Round Robin (RR)

25

Principle

Each process gets the CPU for a limited duration

quantum: generally 10-100 milliseconds

When it has spent its quantum, a process goes back to the Ready queue

Fairness

Consider n processes in the Ready queue, let the quantum value be q

No process ever waits more than $(n-1)q$

Performance tradeoff


- q is large \Rightarrow FIFO
- q is small (close to duration of context switch) \Rightarrow high overhead

Example of RR with $Q = 20$

26

Process	CPU time
P_1	53
P_2	17
P_3	68
P_4	24


Gantt diagram:



Typically induces a higher average turnaround time than SJF, but a smaller response time

Quantum vs. Switch

27



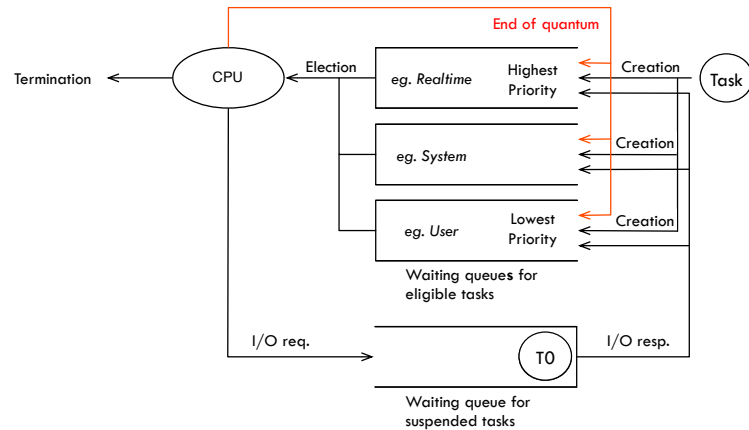
Multilevel Queues

28

- Separate waiting queues
 - foreground (interactive)
 - background (batch)
- Each queue enforces a different scheduling policy
 - foreground – RR
 - background – FCFS
- Requires inter-level scheduling
 - Fixed priority scheduling
 - ie. serve all processes of FG queue, then those of BG queue
 - \Rightarrow May induce famine
 - Time slice
 - Each queue gets a percentage of the available CPU time
 - eg. 80% for FG with RR and 20% for BG with FCFS

Multilevel Queue Scheduling

29



Multilevel Feedback Queues (MFQ)

30

A process can be moved to another queue


Implementation of ageing

Required settings for MFQ

- ▣ # of queues
- ▣ Scheduling policy for each queue
- ▣ Formula for promoting a process to a higher priority queue
- ▣ Formula for demoting a process to a lower priority queue
- ▣ Formula for assigning a queue to a new process

Example of MFQ Scheduling

31



32

Extended Scheduling Policies

Multicore Scheduling

33

- CPU scheduling is more complex
- *Homogeneous processors*
- *Load sharing*
- *Asymmetric Multiprocessing*
 - One processor manages all OS structures
 - No data sharing

Realtime Scheduling

34

- Common misconception: real time does not mean fast!
 - ▣ Real-time processes have timing constraints
 - ▣ Expressed as deadlines or rate requirements
- *Hard Realtime Systems*
 - Guarantee that every process respects deadlines/rates
- *Soft Realtime Systems*
 - May miss deadlines, but QoS then degrades

Common RT Scheduling Policies

35

- Rate monotonic
 - ▣ Just one scalar priority related to the periodicity of the job
 - ▣ $\text{Priority} = 1/\text{rate}$
 - ▣ Static
- Earliest deadline first (EDF)
 - ▣ $\text{Priority} = \text{deadline}$
 - ▣ Dynamic but more complex
- Both require admission control to provide guarantees

Linux Scheduling

36

Two policies: time sharing and real time

Time sharing

- ▣ Priority based on credits
 - processus with the largest amount of credits gets elected
- ▣ Credit subtraction upon every clock tick
- ▣ When credit = 0, election of another process
- ▣ General recreditation when all processes are down to 0 credit
 - Per process recreditation based on priority and history

Real time

- ▣ Soft real time
- ▣ Posix.1b – two classes
 - FCFS and RR
 - Process with highest priority runs first

Linux Scheduling

37

