

# OPERATING SYSTEMS CSCI-SHU215

## Lesson 07 - Signals

## Introduction to signals

2

Most basic inter-process communication mechanism

A signal is information transmitted to a program during its execution

- The system communicates with process users via this mechanism:
  - upon errors (memory violation, I/O error)
  - at the request of the user via the keyboard (ctrl-C, ctrl-Z...)
  - upon disconnection of the line/terminal, etc.
- Given the right (UID) permissions, a process may send a signal to another process
- Upon reception, signal handling induces a default behavior

## Introduction to signals

3

Each signal is associated with a strictly positive integer value that is smaller than **NSIG** (non POSIX constant)

Each value is mapped to a name constant

- `/usr/include/signal.h`
- List of signals:  
`$ kill -l`
- Use the name instead of the number  
Example: **SIGKILL** (= 9), **SIGINT** (= 2), etc.  
`$ kill -KILL <num. proc>; kill -INT <num. proc>`
- Sending a signal comes down to sending an integer value to a process

## Signals - Terminology


4

- Signal delivery  
A signal is delivered to a process when the process takes into account and performs the action associated with it.
- Pending signal  
Signal that was sent to a process but has yet to be delivered.
  - Signal reception is stored in the receiver's PCB
  - "One signal may hide another"  
Reception of a new signal on top of a pending signal with the same value causes the loss of the new signal
- Masked or blocked signal  
Deliberate adjournment of delivery for a signal value

## Behavior upon delivery of a signal

5

Default behaviors are associated with the delivery of each signal value



Do not mix up with interrupts  
Hardware: clock, disk, etc.

## Main POSIX signals

6

<i>Name</i>	<i>Event</i>	<i>Behavior</i>
<b>Termination</b>		
SIGINT	ctrl-C	termination
SIGQUIT	<QUIT> ctrl-\	termination + core
SIGKILL	Kill a process	termination
SIGTERM	Termination signal	termination
SIGCHLD	Terminate or suspend a child process	ignored
SIGABRT	Abnormal termination	termination + core
SIGHUP	Disconnect terminal	termination

## Main POSIX signals

7

<i>Name</i>	<i>Event</i>	<i>Behavior</i>
<b>Suspension / Resumption</b>		
SIGSTOP	Suspend execution	suspension
SIGTSTP	Suspend execution (ctrl-Z)	suspension
SIGCONT	Continuation of a suspended process	resumption
<b>Error notifications</b>		
SIGFPE	Arithmetical error	termination + core
SIGBUS	Bus error	termination + core
SIGILL	Illegal instruction	termination + core
SIGSEGV	Protected memory violation	termination + core
SIGPIPE	Write error on a pipe without reader	termination

## Main POSIX signals

8

<i>Name</i>	<i>Event</i>	<i>Behavior</i>
<b>Other</b>		
SIGALRM	Timeout	termination
SIGUSR1	Reserved for the user	termination
SIGUSR2	Reserved for the user	termination
SIGTRAP	Trace / breakpoint trap	termination + core
SIGIO	Asynchronous I/O	termination

## 9 Sending and Receiving Signals

## Sending a signal

10

### Source code

```
int kill (pid pid_t, int signal)
    pid
    >0      process whose PID is pid
    0       all processes in the same group as the sender
    -1      not specified by POSIX (most systems: broadcast to all processes)
    <-1     all processes in group whose GID is |pid|
    signal
    value between 0 and NSIG (0 sends no signal, tests the existence of pid)
```

### Command line

```
$ kill -l                list all signals
$ kill -sig pid          send signal sig to process pid
```

## Sending a signal - example

11

```
#define _XOPEN_SOURCE 700
#include <sys / types.h>
#include <signal.h>
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>

int main (int arg, char ** argv) {
    printf ("begin program\n");
    / * Send a SIGINT to oneself * /
    kill (getpid(), SIGINT);
    printf ("end program\n");
    return EXIT_SUCCESS;
}
```

## Signal delivery

12

### Delivery occurs when the process **switches from active kernel mode to active user mode**

ie. return from a system call, return from a hardware interrupt, election by the scheduler

### Default behaviors upon reception


- Process termination
- Termination of the process + production of a core file
- Signal ignored
- Suspension of the process (*stopped* or *suspended*)
- Resumption of the process

### Behavior can be redefined by installing a new handler

- **SIG\_IGN** (Ignore the signal)
  - User-defined **handler function**
  - **SIG\_DFL** (Restore the default behavior)
- Applicable to all signals **except** **SIGKILL**, **SIGSTOP**

## User-defined signal delivery

13



## Redefinition of the default handler

14

```
struct sigaction {  
    void (* sa_handler) ();    /* Function */  
    sigset_t sa_mask;          /* Mask signals */  
    int sa_flags;              /* Options */  
};
```

Structure sigaction holds the data that defines signal reception behavior

- *sa\_handler*  
function to execute, SIG\_DFL (default behavior), or SIG\_IGN (discard signal)
- *sa\_mask*  
list of signals added to the current mask during the execution of the handler
  - *sa\_mask U {sig}*  
the value of the signal being delivered is **automatically** masked
- *sa\_flags*  
behavioral options

## Enforcing a new handler

15

```
int sigaction (    int sig,  
                  struct sigaction * act,  
                  struct sigaction * old    );
```

Redefines the handler *act* for the signal *sig*

- *act* and *old* point to a *struct sigaction*
- Delivery of signal *sig* triggers the function *act.sa\_handler*
- *old*: if not NULL, references the previous handler

## Sigaction - example

16

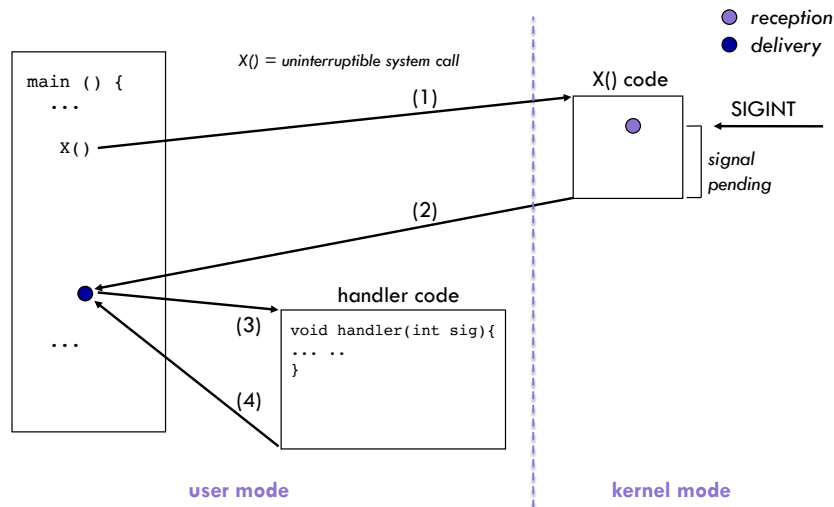
```
void sig_hand(int sig) {  
    printf ( "received signal %d \n", sig);  
}  
  
int main (int argc, char ** argv) {  
    sigset_t sig_proc;  
    struct sigaction action;  
  
    sigemptyset (&sig_proc);  
    action.sa_mask = sig_proc;  
    action.sa_flags = 0;  
    action.sa_handler = sig_hand;  
  
    sigaction (SIGINT, &action, 0);  
  
    kill (getpid (), SIGINT);  
    printf ("end program \n");  
  
    return EXIT_SUCCESS;  
}
```

**sigaction-Ex.C**

> **Sigaction-ex**  
signal received 2  
end program

## Uninterruptible system calls

17



## Interruptible system calls

18

A suspended process whose priority level is interruptible awakens upon receiving a signal

- ▣ Process goes to the ready state
- ▣ The signal gets delivered during the election process  
Execution of the *handler* function
- ▣ Examples of interruptible system calls:
  - `pause`,
  - `sigsuspend`,
  - `wait / waitpid`
  - `read, write`,
  - `etc.`

## Interruptible system calls

19

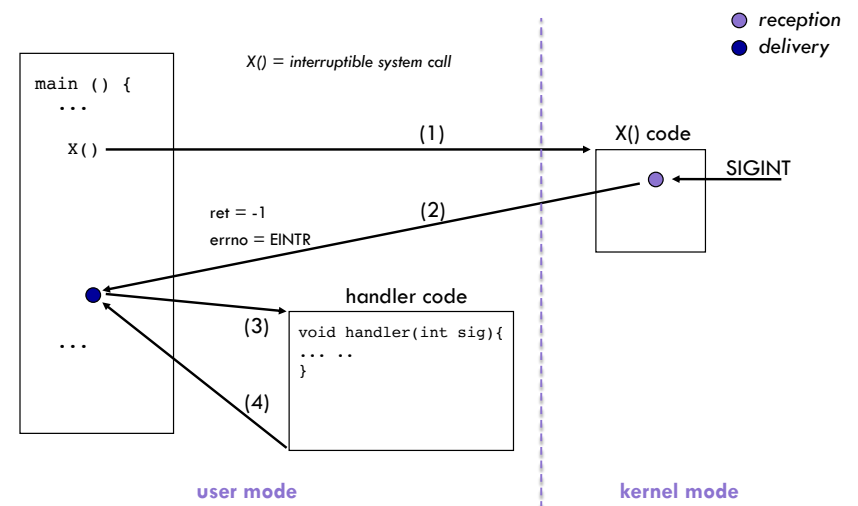
The reception of a signal on an interruptible system call disrupts the call

- ▣ System call returns -1, indicating a failure
- ▣ Error code: `errno = EINTR`.
- ▣ example:
 

```
ret = read (desc, buffer, BUFSIZE)
if ((ret == -1) && (errno == EINTR))
    printf ("signal interrupt\n");
```
- ▣ User-defined handler can trigger automatic recovery

## Interruptible system calls

20



21

## Postponing Signal Delivery

## Postponement of signal delivery

22

A process can choose to postpone the delivery of some signals

Specifies a set of signals whose delivery gets deferred (*mask*)

Cannot defer *SIGKILL* and *SIGSTOP*

```
int sigprocmask (int how, const sigset_t *set, sigset_t *old);
```

<i>how</i>	<i>SIG_BLOCK</i>	union of current mask and set
	<i>SIG_UNBLOCK</i>	unmask set
	<i>SIG_SETMASK</i>	block only signals in set
<i>set</i>	Signal mask	
<i>old</i>	if not NULL, stores value of the previous mask	

Call applies a new mask that replaces the current mask

New mask consists of *set*, or a compound of *set* and the previous mask

## Signal set manipulation

23

The following functions do not change the signals themselves, but manipulate variables that define sets of signals

- ▣ `int sigemptyset (sigset_t *set);`
- ▣ `int sigfillset (sigset_t *set);`
- ▣ `int sigaddset (sigset_t *set, int sig);`
- ▣ `int sigdelset (sigset_t *set, int sig);`
- ▣ `int sigismember (sigset_t *set, int sig);`  
sigismember returns 1 if the signal is a member of the set, 0 otherwise

A process can list blocked signals that are pending

```
int sigpending (sigset_t *set);
```

## Masking signals - example

24

**sigprocmask-ex.c**

```
#define _XOPEN_SOURCE 700
#include <signal.h> // and others
int main (int argc, char ** argv) {
    sigset_t sig_proc;
    printf ("start program \n");

    sigemptyset (&sig_proc);
    sigaddset (&sig_proc, SIGINT);
    sigprocmask (SIG_BLOCK, &sig_proc, NULL);

    sleep (10);

    SIGINT pending
    printf ("after sleep\n");
    sigprocmask (SIG_UNBLOCK, &sig_proc, NULL);
    SIGINT delivered


    printf ("end program \n");

    return EXIT_SUCCESS ;
}
```

> **sigprocmask-ex**  
start program  
> **ctrl-C** ———  
after sleep

## Masking signals

25



## Pending signals – example

26

```
int main (int argc, char * argv []) {  
    sigset_t sig_set; /* List of blocked signals */  
  
    sigemptyset (&sig_set);  
    sigaddset (&sig_set, SIGINT);  
  
    /* Hide SIGINT */  
    sigprocmask (SIG_SETMASK, &sig_set, NULL);  
  
    kill (getpid (), SIGINT);  
  
    /* Get the list of pending signals */  
    sigpending (& sig_set);  
  
    if (sigismember (&sig_set, SIGINT))  
        printf ("SIGINT is pending\n");  
  
    return EXIT_SUCCESS;  
}
```

27

## Waiting for a Signal

## Waiting for a signal

28

A process can suspend its execution voluntarily in order to wait for a signal

- switches to the suspended state
- gets awakened upon reception of an unmasked signal

### **int pause (void)**

- Suspends process until reception of an unmasked signal

### **int sigsuspend (sigset\_t \*p\_ens)**

- Temporarily replaces current mask by *p\_ens*
- If *p\_ens* un.masks pending signals, these get delivered
- Otherwise process is suspended until reception of an unmasked signal
- Reinstates previous mask upon return of the function

## Waiting for a signal – example

29

```
void sig_hand (int sig) {
    printf ( "received signal %d \n", sig);
}

int main (int argc, char ** argv) {
    sigset_t sig_proc;
    struct sigaction action;
    /* Mask SIGINT */
    sigemptyset (&sig_proc);
    sigaddset (&sig_proc, SIGINT);
    sigprocmask (SIG_SETMASK, & sig_proc, NULL);
    /* Redefine handler for SIGINT */
    action.sa_mask = sig_proc;
    action.sa_flags = 0;
    action.sa_handler = sig_hand;
    sigaction (SIGINT, &action, NULL);
    /* Wait for SIGINT */
    sigfillset (&sig_proc);
    sigdelset (&sig_proc, SIGINT);
    sigsuspend (&sig_proc);
    return EXIT_SUCCESS;
}
```

## Timer notification

30

Goal: interrupt the process after a time limit

- Arm a timer. Upon deadline, the process receives a signal.
- Two system calls
  - **alarm** : only real-time, resolution is in seconds, sends SIGALRM
  - **setitimer** : set different types of timers, finer resolution than the second, sends SIGALRM, SIGTVALRM, or SIGPROF
- Default behavior is process termination

## Timer notification

31

`unsigned alarm (unsigned seconds);`

- Timer in seconds
  - Real time (wall-clock time)
  - Resolution is the second
- Sends SIGALRM upon expiration
- One call at a time per process
  - A new call arms a new timer and cancels the previous one
  - Calling with value 0 does not arm a new timer
- Return value
  - Amount of time left on the timer from a previous call
  - 0 if no alarm is currently set

## Timer notification – example

32

```
void sig_hand (int sig) {
    printf ("received signal %d \n", sig);
    alarm (1);
}
```

**sig\_ALARM.c**

```
int main (int argc, char ** argv) {
    sigset_t sig_proc; struct sigaction share;
    sigemptyset (& sig_proc);
    action.sa_mask = sig_proc;
    action.sa_flags = 0;
    action.sa_handler = sig_hand;
    sigaction (SIGALRM, & action, 0);
    while (1) {
        alarm(1);
        pause ();
    }
    return EXIT_SUCCESS;
}
```

```
> sig_ALARM
received signal 20
received signal 20
received signal 20
... ..
```



33

## Tricky Situations

## Signals vs. Parenthood

34

Signal SIGCHLD notifies a parent process about changes of state incurred by its children

Automatically sent upon

- ▣ child termination
- ▣ child suspension: child receives SIGSTOP or SIGTSTP.

The default behavior is to ignore SIGCHLD

## Signals vs Inheritance

35

A child process

- ▣ does not inherit the parent's pending signals
- ▣ inherits the parent's signal mask and handlers

`exec()` resets the default handlers

## Loss of pending signals

36

"One signal may hide another"

The reception of a signal sets this signal value as pending

The reception of a new signal on top of a pending signal ***never gets delivered***

## Loss of pending signals – example

37

```
int cont;

void sig_hand (int sig) {
    if (sig == SIGUSR1)
        cont ++;
    else {
        printf ("nb of SIGUSR1 received: %d \n", cont);
        exit (0);
    }
}
```

sig\_contUSR1.c

## Loss of pending signals – example

38

```
pid_t pid_child;

int main (int argc, char ** argv) {
    sigset_t sig_proc; int i; struct sigaction action;
    sigemptyset (& sig_proc);
    action.sa_mask = sig_proc;
    action.sa_flags = 0;
    action.sa_handler = sig_hand;
    sigaction (SIGUSR1, & action, 0);
    sigaction (SIGINT, & action, 0);
    if ((pid_child = fork ()) == 0) {
        while (1)
            pause ();
    } else {
        for (i = 0; i <20; i ++)
            kill (pid_child, SIGUSR1);
        kill (pid_child, SIGINT);
        wait (NULL);
        return EXIT_SUCCESS;
    }
}
```

sig\_contUSR1.c

> **sig\_contUSR1**  
nb of SIGUSR1 received: 4

## Limitations of signals

39

A signal notifies an event, but provides little more information than the type of event (signal value)

In particular, POSIX.1 signals do not allow to know

- ▣ the number of receptions of a signal value in-between deliveries
- ▣ the date of reception of a signal
  - Most systems handle deliveries by order of increasing signal value
- ▣ the PID of the process that sent the signal